

Salvador Fco Criado Melero

Senior AI Engineer — LLM Applications, Agentic Systems & Model Deployment

Hands-on senior engineer with **8+ years** owning distributed systems in production, now focused on *applied AI end-to-end*. I build the **application layer** — RAG over private corpora, *multi-agent orchestration with tool calling*, real-time voice agents under 2 s, guardrails and evaluation — and the **infrastructure underneath**: *quantisation with LLM Compressor, serving on vLLM and Triton, benchmarking with GuideLLM and lm-eval, Langfuse tracing and golden datasets*. Backbone is classic engineering: deep AWS, expert Terraform, Kubernetes and event-driven systems at scale.



Granada, Spain | +34 635 656 650 | crimelero@hotmail.com | linkedin.com/in/salvador-cm

§02 SKILLS

LLM Applications & Agents

RAG · Embeddings
Multi-agent · Tool calling
Vector DBs · Qdrant
Real-time voice (STT+LLM+TTS)
Structured output · Guardrails · PII
OCR · Document AI · Human-in-the-loop
LangChain · LlamaIndex · CrewAI

Model Deployment & Evaluation

vLLM · NVIDIA Triton
Quantisation · LLM Compressor
Langfuse · Golden sets
GuideLLM · lm-eval
Mistral · Llama (self-hosted)
AWS Bedrock · Claude · GPU inference

Software Engineering

Python · Node.js · NestJS · Django
REST · GraphQL · API design
Microservices · Event-driven · Time-series
PostgreSQL · MongoDB · ES

Cloud & DevOps

AWS (deep) · Terraform
Kubernetes · Helm · Docker
Azure · AKS
Serverless · Linux · Bash
CI/CD · Observability

§03 CERTIFICATIONS

- AWS Certified DevOps Engineer — Professional (2022)
- Agentic AI — DeepLearning.AI, Andrew Ng (2026)
- Certified Scrum Master — CSM, European Scrum (2023)

§04 EDUCATION

MSc Electronic Systems

UPM Madrid · 2017–2018

BSc Telecommunications Eng. — Electronics

University of Granada (UGR) · 2013–2017

§05 LANGUAGES

Spanish **NATIVE**
English **PROFESSIONAL**
French **BASIC**

§01 EXPERIENCE

Independent Consultant · AI & Platform Engineering

Sep 2025 — Present

SELECTED CLIENT ENGAGEMENTS · REMOTE

Enterprise Contact-Centre Client — On-prem agentic voice platform

Sep 2025 — Jun 2026

- End-to-end ownership of an agentic LLM product on the customer's own GPU** — strict on-prem: data never leaves the network. Agentic orchestration of inbound and outbound flows with tool calling against the client's CRM, RAG on Qdrant grounding the agent in the client's own product and procedure documents, and a streaming voice loop (STT + LLM + TTS) under **2 s** end-to-end.
- Model deployment, optimisation and benchmarking**. Serving on vLLM + Mistral, quantisation with LLM Compressor to fit memory and latency on the available GPU, throughput benchmarking with GuideLLM and lm-evaluation-harness.
- Evaluation and observability as first-class features**. Per-turn Langfuse tracing, golden-dataset regression and structured logging — quality measured, failures diagnosable without re-running the call. *Stack: Python, vLLM + Mistral, LLM Compressor, Qdrant, GuideLLM, Langfuse, Docker.*

Document Processing Automation — OCR + LLM pipeline

Jan 2026 — Present

- Bulk document intake turned into structured, postable data**. Invoices and receipts ingested at volume, document type and target account classified by an LLM over OCR text, accounting entries emitted in seconds, with human-in-the-loop review of exceptions and an audit trail per entry — on AWS Bedrock + Claude with Langfuse tracing over a Python / Django backend, plus structured output, PII guardrails and golden-set regression.

Westinghouse — Full AWS → Azure migration

Mar 2026 — Present

- Multi-tenant e-learning platform moved end-to-end from AWS to Azure** — **100k+** users, workloads onto AKS (Kubernetes), infrastructure as code in Terraform — a returning client, on a platform I originally helped build. *Stack: Azure, AKS, Terraform, Docker, PHP / Laravel, Moodle.*

Lead Software Engineer at Iyris

Aug 2024 — Dec 2025

AWS DATA & AI PLATFORM · UNITED KINGDOM · REMOTE

- Owned the product's AI layer over five data**. LLM assistant querying and operating on real-time telemetry for end users.
- Platform ownership after downsize, plus support for Data Science**. Took over a platform previously run by **20+** engineers, leading a small crew through redesign, and built what the DS team needed: data lake on S3 + Glue + Athena, BI tooling and processing pipelines.
- High-throughput event-driven ingestion, fully in Terraform**. Telemetry from **1,000+** sensors as time-series with critical-event detection and alert routing, on Lambda, IoT Core, S3 and SQS. *Stack: AWS, Python, Node.js, Terraform, MQTT, time-series.*

Principal Software Engineer at TMC Spain

Jan 2022 — Aug 2024

CONSULTING FOR WESTINGHOUSE · MADRID · REMOTE

- Technical lead of a multi-tenant e-learning platform, running the engineering team**. Moodle-based platform with **80k+** users per tenant on PHP / Laravel / Vue with Elasticsearch search — technical direction, delivery and production support.
- Owned the infrastructure and DevOps side**. Whole multi-tenant estate moved to Terraform on AWS and Azure with automated deployments on Azure DevOps — deploy time cut from **1 week** to **1 hour** — plus Docker / Kubernetes workloads and CI/CD.
- Real-time alert filtering for nuclear plants**. MQTT detection pipeline on NestJS, **1,000+** signals/second, where accuracy and latency are non-negotiable.

Full-Stack Engineer at Nazaries IT

Feb 2018 — Jan 2022

IOT PLATFORM · HARDWARE TO APPLICATION · GRANADA

- Built and maintained an IoT platform of 1,000+ field sensors across the whole stack** — hardware, firmware, backend and application: sensor electronics and PCB in Altium, embedded firmware in C and MicroPython, a new BLE / BLE Mesh sensor line, telemetry ingestion and time-series services, and the customer-facing web and mobile app on top.
- Monolith to microservices on the data-critical patfis**. Moved from Ruby on Rails to Node.js / NestJS — queries previously impossible (10 sensors, 1-minute resolution, 10 years) returning under **30 s**. *Stack: Node.js, NestJS, Rails, MQTT, MongoDB, BLE Mesh, C.*